

# PracticeSession07

June 11, 2025

## Numerical Linear Algebra for Computational Science and Information Engineering

### Orthogonalization and Least-Squares Problems

Nicolas Venkovic (nicolas.venkovic@tum.edu)

```
[1]: using LinearAlgebra, Plots, Printf, Latexify, LaTeXStrings, BenchmarkTools, Random, SparseArrays
```

```
[2]: function conditioned_matrix(m::Int, n::Int, cond::Float64)
    U, _ = qr(randn(m, m))
    V, _ = qr(randn(n, n))
    s = range(1.0, 1/cond; length=n)
    Σ = Diagonal(s)
    A = U[:, 1:n] * Σ * V'
    return A
end;
```

### Exercise #1: Implement and test your implementation of HouseholderQR

```
[3]: function HouseholderQR(A)
    m, n = size(A)
    Q = Matrix{Float64}(I, m, m)
    R = copy(A)
    for k = 1:n
        # Compute Householder vector
        v = R[k:m, k]
        v[1] -= sign(v[1]) * norm(v) # usage of sign(v[1]) allows to prevent
        ↪catastrophic cancellation
        beta = 2. / v'v
        # Apply Householder reflection
        R[k:m, k:n] -= beta .* v * (v'R[k:m, k:n])
        Q[1:m, k:m] -= Q[1:m, k:m] * v * (beta .* v')
    end
    return Q, R
end;
```

```
[4]: m = 1_000;
      n = 200;
      Random.seed!(3);
      A1 = conditioned_matrix(m, n, 1e8);
      A2 = conditioned_matrix(m, n, 1e10);
```

```
[5]: Q, R = @btime HouseholderQR(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)

      2.437 s (7807 allocations: 5.94 GiB)
      ||A - QR|| = 1.15E-14, ||I - Q'Q|| = 4.62E-14
```

```
[6]: Q, R = @btime HouseholderQR(A2);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)

      2.519 s (7807 allocations: 5.94 GiB)
      ||A - QR|| = 1.09E-14, ||I - Q'Q|| = 4.54E-14
```

## Exercise #2: Implement CholeskyQR, CholeskyQR2 and Shifted CholeskyQR3

```
[7]: function CholeskyQR(A)
      X = A'A
      chol = cholesky(X)
      Q = A / chol.U
      return Q, chol.U
    end;
```

```
[8]: function CholeskyQR2(A)
      Q1, R1 = CholeskyQR(A)
      Q, R2 = CholeskyQR(Q1)
      return Q, R2 * R1
    end;
```

```
[9]: function ShiftedCholeskyQR(A)
      X = A'A
      s = 11 * (m * n + n * (n + 1)) * eps() * norm(A)
      X += s * I
      chol = cholesky(X)
      Q = A / chol.U
      return Q, chol.U
    end;
```

```
[10]: function ShiftedCholeskyQR3(A)
      Q1, R1 = ShiftedCholeskyQR(A)
      Q, R2 = CholeskyQR2(Q1)
      return Q, R2 * R1
    end;
```

### Exercise #3: Compare loss of orthogonality and residual of each CholeskyQR implementation

```
[11]: Q, R = @btime CholeskyQR(A1);  
@printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)
```

```
2.357 ms (10 allocations: 2.14 MiB)  
||A - QR|| = 9.30E-16, ||I - Q'Q|| = 2.77E-02
```

```
[12]: Q, R = @btime CholeskyQR(A2);  
@printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)
```

PosDefException: matrix is not positive definite; Factorization failed.

Stacktrace:

```
[1] checkpositivedefinite  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/factorization.jl:68 [inlined]  
[2] #cholesky!#163  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:269 [inlined]  
[3] cholesky!  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:267 [inlined]  
[4] #cholesky!#164  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:301 [inlined]  
[5] cholesky! (repeats 2 times)  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:295 [inlined]  
[6] _cholesky  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:411 [inlined]  
[7] cholesky(A::Matrix{Float64}, ::NoPivot; check::Bool)  
  @ LinearAlgebra ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/  
↳ stdlib/v1.11/LinearAlgebra/src/cholesky.jl:401  
[8] cholesky (repeats 2 times)  
  @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/  
↳ LinearAlgebra/src/cholesky.jl:401 [inlined]  
[9] CholeskyQR(A::Matrix{Float64})  
  @ Main ./In[7]:3  
[10] var"##core#256"()  
  @ Main ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:598  
[11] var"##sample#257"(:::Tuple{}, __params::BenchmarkTools.Parameters)  
  @ Main ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:607  
[12] _lineartrial(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters;   
↳ maxevals::Int64, kwargs::@Kwargs{ })  
  @ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:18
```

```

[13] _lineartrial(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters)
      @ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:18
[14] #invokelatest#2
      @ ./essentials.jl:1055 [inlined]
[15] invokelatest
      @ ./essentials.jl:1052 [inlined]
[16] #lineartrial#46
      @ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:51 [inlined]
[17] lineartrial
      @ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:50 [inlined]
[18] tune!(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters;
      ↪progressid::Nothing, nleaves::Float64, ndone::Float64, verbose::Bool, pad::
      ↪String, kwargs::@Kwargs{})
      @ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:29
[19] tune!
      @ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:288 [inlined]
[20] tune!(b::BenchmarkTools.Benchmark)
      @ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:28
[21] top-level scope
      @ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:728

```

```

[13]: Q, R = @btime CholeskyQR2(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)

      4.780 ms (22 allocations: 4.58 MiB)
      ||A - QR|| = 3.35E-15, ||I - Q'Q|| = 5.57E-15

```

```

[14]: Q, R = @btime CholeskyQR2(A2);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)

```

PosDefException: matrix is not positive definite; Factorization failed.

Stacktrace:

```

[1] checkpositivedefinite
      @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
      ↪LinearAlgebra/src/factorization.jl:68 [inlined]
[2] #cholesky!#163
      @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
      ↪LinearAlgebra/src/cholesky.jl:269 [inlined]
[3] cholesky!
      @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
      ↪LinearAlgebra/src/cholesky.jl:267 [inlined]
[4] #cholesky!#164
      @ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
      ↪LinearAlgebra/src/cholesky.jl:301 [inlined]
[5] cholesky! (repeats 2 times)

```

```

@ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
↳LinearAlgebra/src/cholesky.jl:295 [inlined]
[6] _cholesky
@ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
↳LinearAlgebra/src/cholesky.jl:411 [inlined]
[7] cholesky(A::Matrix{Float64}, ::NoPivot; check::Bool)
@ LinearAlgebra ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/
↳stdlib/v1.11/LinearAlgebra/src/cholesky.jl:401
[8] cholesky (repeats 2 times)
@ ~/.julia/juliaup/julia-1.11.5+0.x64.linux.gnu/share/julia/stdlib/v1.11/
↳LinearAlgebra/src/cholesky.jl:401 [inlined]
[9] CholeskyQR(A::Matrix{Float64})
@ Main ./In[7]:3
[10] CholeskyQR2(A::Matrix{Float64})
@ Main ./In[8]:2
[11] var"##core#270"()
@ Main ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:598
[12] var"##sample#271"(:::Tuple{}, __params::BenchmarkTools.Parameters)
@ Main ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:607
[13] _lineartrial(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters;↳
↳maxevals::Int64, kwargs::@Kwargs{})
@ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:18
[14] _lineartrial(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters)
@ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:18
[15] #invokelatest#2
@ ./essentials.jl:1055 [inlined]
[16] invokelatest
@ ./essentials.jl:1052 [inlined]
[17] #lineartrial#46
@ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:51 [inlined]
[18] lineartrial
@ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:50 [inlined]
[19] tune!(b::BenchmarkTools.Benchmark, p::BenchmarkTools.Parameters;↳
↳progressid::Nothing, nleaves::Float64, ndone::Float64, verbose::Bool, pad::
↳String, kwargs::@Kwargs{})
@ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:29
[20] tune!
@ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:288 [inlined]
[21] tune!(b::BenchmarkTools.Benchmark)
@ BenchmarkTools ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:28
[22] top-level scope
@ ~/.julia/packages/BenchmarkTools/1i1mY/src/execution.jl:728

```

```

[15]: Q, R = @btime ShiftedCholeskyQR3(A1);
@printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)

```

7.750 ms (51 allocations: 7.33 MiB)

```
||A - QR|| = 3.89E-15, ||I - Q'Q|| = 4.35E-15
```

```
[16]: Q, R = @btime ShiftedCholeskyQR3(A2);
@printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)
```

```
8.398 ms (51 allocations: 7.33 MiB)
||A - QR|| = 3.93E-15, ||I - Q'Q|| = 4.27E-15
```

#### Exercise #4: Implement CGS, MGS and CGS2

```
[17]: function CGS_naive(X::Array{Float64,2})
    n, m = size(X)
    Q = Array{Float64,2}(undef, n, m)
    R = zeros(Float64, m, m)
    R[1, 1] = norm(X[:, 1])
    Q[:, 1] = X[:, 1] ./ R[1, 1]
    for i in 2:m
        Q[:, i] = X[:, i]
        R[1:i-1, i] = Q[:, 1:i-1]'X[:, i] # creates temporary array => very slow
        Q[:, i] -= Q[:, 1:i-1] * R[1:i-1, i] # creates temporary array => very slow
        R[i, i] = norm(Q[:, i])
        Q[:, i] ./= R[i, i]
    end
    return Q, R
end;
```

```
function CGS(X::Array{Float64,2})
    n, m = size(X)
    Q = zeros(Float64, n, m)
    R = zeros(Float64, m, m)
    r = zeros(Float64, m)
    w = zeros(Float64, n)
    R[1, 1] = norm(X[:, 1])
    Q[:, 1] = X[:, 1] ./ R[1, 1]
    for i in 2:m
        w = X[:, i]
        r = (w'Q)'
        r[i:m] = 0.
        R[:, i] = r
        w -= Q * r
        R[i, i] = norm(w)
        Q[:, i] = w ./ R[i, i]
    end
    return Q, R
end;
```

```
[18]: function MGS_naive(X::Array{Float64,2})
    n, m = size(X)
```

```

Q = Array{Float64,2}(undef, n, m)
R = zeros(Float64, m, m)
R[1, 1] = norm(X[:, 1])
Q[:, 1] = X[:, 1] ./ R[1, 1]
for i in 2:m
    Q[:, i] = X[:, i]
    for j in 1:i-1
        R[j, i] = Q[:, j]'Q[:, i]
        Q[:, i] -= R[j, i] .* Q[:, j]
    end
    R[i, i] = norm(Q[:, i])
    Q[:, i] ./= R[i, i]
end
return Q, R
end;

function MGS(X::Array{Float64,2})
    n, m = size(X)
    Q = Array{Float64,2}(undef, n, m)
    q = Vector{Float64}(undef, n)
    R = zeros(Float64, m, m)
    w = Vector{Float64}(undef, n)
    R[1, 1] = norm(X[:, 1])
    Q[:, 1] = X[:, 1] ./ R[1, 1]
    for i in 2:m
        w .= X[:, i]
        for j in 1:i-1
            q .= Q[:, j]
            R[j, i] = q'w
            w -= R[j, i] .* q
        end
        R[i, i] = norm(w)
        Q[:, i] = w ./ R[i, i]
    end
    return Q, R
end;

```

```

[19]: function CGS2_naive(X::Array{Float64,2})
    n, m = size(X)
    Q = Array{Float64,2}(undef, n, m)
    R = zeros(Float64, m, m)
    R[1, 1] = norm(X[:, 1])
    Q[:, 1] = X[:, 1] ./ R[1, 1]
    for i in 2:m
        Q[:, i] = X[:, i]
        R[1:i-1, i] .= Q[:, 1:i-1]'X[:, i] # creates temporary array => very slow
        Q[:, i] -= Q[:, 1:i-1] * R[1:i-1, i] # creates temporary array => very slow
    end
end;

```

```

    r = Q[:, 1:i-1]'Q[:, i] # creates temporary array => very slow
    Q[:, i] -= Q[:, 1:i-1] * r # creates temporary array => very slow
    R[i, i] = norm(Q[:, i])
    Q[:, i] ./= R[i, i]
end
return Q, R
end;

function CGS2(X::Array{Float64,2})
    n, m = size(X)
    Q = zeros(Float64, n, m)
    R = zeros(Float64, m, m)
    r = zeros(Float64, m)
    w = zeros(Float64, n)
    R[1, 1] = norm(X[:, 1])
    Q[:, 1] = X[:, 1] ./ R[1, 1]
    for i in 2:m
        w .= X[:, i]
        r .= Q'w
        r[i:m] .= 0.
        R[:, i] = r
        w -= Q * r
        r .= Q'w
        r[i:m] .= 0.
        w -= Q * r
        R[i, i] = norm(w)
        Q[:, i] = w ./ R[i, i]
    end
    return Q, R
end;

```

**Exercise #5: Compare runtime, loss of orthogonality and residual of CGS, MGS and CGS2**

```

[20]: Q, R = @btime CGS_naive(A1);
      Q, R = @btime CGS(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)

```

```

      55.181 ms (5985 allocations: 316.59 MiB)
      14.985 ms (2806 allocations: 8.31 MiB)
      ||A - QR|| = 2.45E-15, ||I - Q'Q|| = 9.54E-08

```

```

[21]: Q, R = @btime CGS(A2);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)

```

```

      14.345 ms (2806 allocations: 8.31 MiB)
      ||A - QR|| = 2.47E-15, ||I - Q'Q|| = 1.43E-05

```



```
[22]: Q, R = @btime MGS_naive(A1);
      Q, R = @btime MGS(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)
```

```
140.893 ms (359807 allocations: 925.59 MiB)
82.462 ms (180315 allocations: 464.51 MiB)
||A - QR|| = 3.81E-15, ||I - Q'Q|| = 1.35E-08
```

```
[23]: Q, R = @btime MGS(A2);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)
```

```
73.142 ms (180315 allocations: 464.51 MiB)
||A - QR|| = 3.85E-15, ||I - Q'Q|| = 1.47E-06
```

```
[24]: Q, R = @btime CGS2_naive(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)
      Q, R = @btime CGS2(A1);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A1 - Q * R) norm(I - Q'Q)
```

```
110.824 ms (9965 allocations: 626.57 MiB)
||A - QR|| = 3.30E-15, ||I - Q'Q|| = 5.73E-15 27.846 ms (4398 allocations:
11.69 MiB)
||A - QR|| = 3.30E-15, ||I - Q'Q|| = 5.55E-15
```

```
[25]: Q, R = @btime CGS2(A2);
      @printf "||A - QR|| = %.2E, ||I - Q'Q|| = %.2E" norm(A2 - Q * R) norm(I - Q'Q)
```

```
28.367 ms (4398 allocations: 11.69 MiB)
||A - QR|| = 3.33E-15, ||I - Q'Q|| = 5.36E-15
```

## Exercise #6: Implement Golub-Kahan bidiagonalization and verify matrix transformation

```
[26]: function bidiagonalization(A, r0, k)
      m, n = size(A)
      U = zeros(Float64, m, k + 1)
      V = zeros(Float64, n, k + 1)
      beta = zeros(Float64, k + 1)
      alpha = zeros(Float64, k + 1)
      U[:, 1] = r0
      beta[1] = norm(U[:, 1])
      U[:, 1] = U[:, 1] ./ beta[1]
      V[:, 1] = A'U[:, 1]
      alpha[1] = norm(V[:, 1])
      V[:, 1] = V[:, 1] ./ alpha[1]
      for i = 1:k
          U[:, i + 1] = A * V[:, i] - alpha[i] * U[:, i]
          beta[i + 1] = norm(U[:, i + 1])
          U[:, i + 1] = U[:, i + 1] / beta[i + 1]
```

```

    V[:, i + 1] = A'U[:, i + 1] - beta[i + 1] * V[:, i]
    alpha[i + 1] = norm(V[:, i + 1])
    V[:, i + 1] = V[:, i + 1] / alpha[i + 1]
end
return U, V[:, 1:k], alpha[1:k], beta
end;

```

```

[27]: Random.seed!(3)
m = 100_000; n = 200;
A = sprand(Float64, m, n, .4);
b = rand(m);

```

```

[28]: k = 5;
U, V, alpha, beta = bidiagonalization(A, rand(m), k);
B = spdiags(k + 1, k, 0=>alpha, -1=>beta[2:k+1]);
norm(U'A*V - B)

```

[28]: 6.415087381394998

```

[29]: U'A*V

```

```

[29]: 6×5 Matrix{Float64}:
 773.996      8.41039e-11  1.57414e-7  0.000298512  0.589457
 457.727      4.65152     9.30947e-8  0.00017654  0.348604
 9.09368e-10  96.5391      4.41971    9.91893e-8  0.000195861
 1.70404e-6   4.54171e-9   96.5037    4.36742     1.04521e-7
 0.00322698   8.59641e-6   4.58986e-9 96.6373     4.1883
 6.37841      0.0169915    9.07075e-6 4.78334e-9  96.5428

```

```

[30]: B

```

```

[30]: 6×5 SparseMatrixCSC{Float64, Int64} with 10 stored entries:
 773.996
 457.727  4.65152
          96.5391  4.41971
                96.5037  4.36742
                    96.6373  4.1883
                        96.5428

```

```

[31]: norm(I - U'U)

```

[31]: 0.01003147391463673

```

[32]: norm(I - V'V)

```

[32]: 0.0010770355772659224

## Exercise #7: Implement and test LSQR

```
[33]: function LSQR(A, b, x0, k, tol)
    A_F = sqrt(sum(A.nzval.^2))
    m, n = size(A)
    U = zeros(Float64, m, k + 1)
    V = zeros(Float64, n, k)
    B = spdiags(k + 1, k, 0=>zeros(Float64, k), -1=>zeros(Float64, k))
    U[:, 1] = b - A * x0
    beta1 = norm(U[:, 1])
    U[:, 1] = U[:, 1] ./ beta1
    V[:, 1] = A'U[:, 1]
    B[1, 1] = norm(V[:, 1])
    V[:, 1] = V[:, 1] / B[1, 1]
    beta1_x_e1 = zeros(Float64, k + 1); beta1_x_e1[1] = beta1
    i = 0
    err = 1.
    yi = 1.
    while (err > tol && i < k)
        i += 1
        U[:, i + 1] = A * V[:, i] - B[i, i] * U[:, i]
        B[i + 1, i] = norm(U[:, i + 1])
        U[:, i + 1] = U[:, i + 1] / B[i + 1, i]
        V[:, i + 1] = A'U[:, i + 1] - B[i + 1, i] * V[:, i]
        B[i + 1, i + 1] = norm(V[:, i + 1])
        V[:, i + 1] = V[:, i + 1] / B[i + 1, i + 1]
        yi = B[1:i+1, 1:i] \ beta1_x_e1[1:i+1]
        ti = beta1_x_e1[1:i+1] - B[1:i+1, 1:i] * yi # ||ti|| = ||ri||
        si = B[1:i+1, 1:i+1]'ti # ||si|| = ||A'ri||
        err = norm(si) / (A_F * norm(ti))
        @printf "i = %i, ||A'ri||/(||A||.*||ri||) = %.2E\n" i err
    end
    return x0 + V[:, 1:i] * yi
end;
```

```
[34]: x = LSQR(A, b, rand(n), 100, 1e-5);
norm(b - A * x)
```

```
i = 1, ||A'ri||/(||A||.*||ri||) = 5.75E-02
i = 2, ||A'ri||/(||A||.*||ri||) = 1.07E-02
i = 3, ||A'ri||/(||A||.*||ri||) = 4.53E-04
i = 4, ||A'ri||/(||A||.*||ri||) = 2.17E-05
i = 5, ||A'ri||/(||A||.*||ri||) = 1.07E-06
```

```
[34]: 92.84273159426289
```

```
[35]: x = A \ b;
norm(A * x - b)
```

[35] : 92.84273158157164