

Numerical Linear Algebra for Computer Science and Information Engineering

Lecture 02 Essentials of the Julia Language

Nicolas Venkovic
nicolas.venkovic@tum.de

Group of Computational Mathematics
School of Computation, Information and Technology
Technical University of Munich

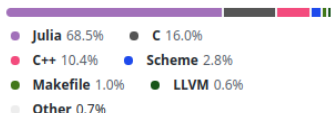
Winter 2025-26



Fact sheet of the Julia language

- ▶ Started at **MIT** in 2009 to develop a **fast open source** and **free high-level** language
- ▶ Features
 - **Dynamically typed** (also enables **static types for better performance**)
 - Just-in-time (JIT) compiled (i.e., **compiled at runtime**)
 - Provided with full-featured interactive command-line **REPL** (read-eval-print loop)
 - Designed for **parallelism** and **distributed computing** (part of the standard library)
 - **No need to vectorize code for performance**
 - Supports notebooks
- ▶ Version 1.0 released in 2018
- ▶ Current release is **1.12.1**
- ▶ Used at MIT, Stanford, UC Berkeley, Amazon, Apple, Google, IBM, Intel, Microsoft, ...
- ▶ Over 45 millions downloads as of January 2023

What is Julia made of?



- ▶ Most of the Julia **standard library** is written in **Julia**
- ▶ Julia makes use of **pre-existing libraries** (mostly in C/C++) for:
 - **BLAS/LAPACK**, however, native Julia versions exist for most functionalities. Optimized native Julia BLAS can match the performances of Intel MKL and OpenBLAS.
 - **Regular expressions** (PCRE)
 - **Downloading files** (libcurl)
 - **Low-level asynchronous IO** (libuv)
 - **Compilation** (LLVM)
 - **Extended precision arithmetic** (GMP, MPFR), but native Julia solutions also exist.

RBGS: An algorithm for comparison purposes

- Randomized block Gram-Schmidt (RBGS) procedure by Balabanov and Grigori (2021):

$$\text{RBGS} : (X, \Theta) \in \mathbb{R}^{n \times m} \times \mathbb{R}^{k \times m} \mapsto (Q, R) \in \mathbb{R}^{n \times m} \times \mathbb{R}^{m \times m}$$

such that $X = QR$ where $m < k \ll n$, $(\Theta Q)^T \Theta Q = I_m$ and $\text{Ran}(\Theta Q) = \text{Ran}(\Theta X)$. We exploit the following structure of p blocks of size $n \times s$:

$$\begin{aligned} X &= [X_{:,1:s}, X_{:,s+1:2s}, \dots, X_{:,(p-1)s+1:ps}] \\ Q &= [Q_{:,1:s}, Q_{:,s+1:2s}, \dots, Q_{:,(p-1)s+1:ps}] \end{aligned}$$

where $m = ps$.

- Exploiting those block structures, we have $(\Theta Q)^T \Theta Q = I_m \implies$

$$R_{(i-1)s+1:is, (j-1)s+1:js} = (\Theta Q_{:,(i-1)s+1:is})^T \Theta X_{:,(j-1)s+1:js}, \quad (i, j) \in [1, p]^2$$

RBGS: An algorithm for comparison purposes

- There are different possible implementations of RBGS algorithm. Let us consider the following:

Algorithm 1 RBGS : $(X, \Theta) \mapsto (Q, R)$

- 1: $\text{RGS}(X_{:,1:s}) \mapsto (Q_{:,1:s}, R_{1:s,1:s}, S_{:,1:s})$ ▷ $S := \Theta Q$
 - 2: **for** $i = 2, \dots, p$ **do**
 - 3: $P := \Theta X_{:,(i-1)s+1:is}$ ▷ Sketching
 - 4: $R_{1:(i-1)s,(i-1)s+1:is} := S_{:,1:(i-1)s}^\dagger P$ ▷ Block least-squares problem
 - 5: $Q_{:,(i-1)s+1:is} := X_{:,(i-1)s+1:is} - Q_{:,1:(i-1)s} R_{1:(i-1)s,(i-1)s+1:is}$ ▷ BLAS-3
 - 6: $\text{RGS}(Q_{:,(i-1)s+1:is}) \mapsto (Q_{:,(i-1)s+1:is}, R_{(i-1)s+1:is,(i-1)s+1:is}, S_{:,(i-1)s+1:is})$
-

where RGS corresponds to RBGS with $s = 1$.

- In what follows, line 3 will be done more efficiently using a matrix-free fast transform.

Julia is close to math

```
function RBGS(X::Array{Float64,2}, p::Int, k::Int)
    n, m = size(X)
    s = Int(m / p)
    P = Array{Float64}(undef, k, s)
    Q = Array{Float64,2}(undef, n, m)
    R = zeros(Float64, m, m)
    S = Array{Float64,2}(undef, k, m)
    srht = set_srht(n, k)
    Q[:, 1:s], R[1:s, 1:s], S[:, 1:s] = RGS(X[:, 1:s], srht)
    for i in 2:p
        P .= MatrixFreeTheta(X[:, (i-1)*s+1:i*s], srht)
        R[1:(i-1)*s, (i-1)*s+1:i*s] = S[:, 1:(i-1)*s] \ P
        Q[:, (i-1)*s+1:i*s] = X[:, (i-1)*s+1:i*s]
            .- Q[:, 1:(i-1)*s] * R[1:(i-1)*s, (i-1)*s+1:i*s]
        Q[:, (i-1)*s+1:i*s],
        R[(i-1)*s+1:i*s, (i-1)*s+1:i*s],
        S[:, (i-1)*s+1:i*s] = RGS(Q[:, (i-1)*s+1:i*s], srht)
    end
    return Q, R, S
end
```

So is Python

```
def RBGS(X, p, k):
    n, m = X.shape
    s = int(m / p)
    P = np.zeros((k, s))
    Q = np.zeros((n, m))
    R = np.zeros((m, m))
    S = np.zeros((k, m))
    srht = set_srht(n, k)
    Q[:, :s], R[:, s, :s], S[:, :s] = RGS(X[:, :s], srht)
    for i in range(1, p):
        P[:, :] = MatrixFreeTheta(X[:, i*s:(i+1)*s], srht)
        R[:, i*s, i*s:(i+1)*s] = np.linalg.lstsq(S[:, :i*s], P)[0]
        Q[:, i*s:(i+1)*s] = X[:, i*s:(i+1)*s]
            - np.matmul(Q[:, :i*s], R[:, i*s, i*s:(i+1)*s])
        Q[:, i*s:(i+1)*s],
        R[i*s:(i+1)*s, i*s:(i+1)*s],
        S[:, i*s:(i+1)*s] = RGS(Q[:, i*s:(i+1)*s], srht)
    return Q, R, S
```

But not C

```
void RBGS(int n, int m, int p, int k, double *X, struct SRHT srht, double *Q, double *R, double *S) {
    int s = m / p;
    double *P = (double*)malloc(k * s * sizeof(double));
    double *Rtmp = (double*)malloc(m * s * sizeof(double));
    double *StS = (double*)malloc(m * m * sizeof(double));
    lapack_int *ipiv = (lapack_int*)malloc(m * sizeof(lapack_int));
    RGS(n, s, k, &X[0], srht, &Q[0], Rtmp, &S[0]);
    for (int v=0; v<s; v++)
        for (int u=0; u<s; u++)
            R[v * m + u] = Rtmp[v * s + u];
    for (int v=0; v<s; v++)
        for (int u=s; u<m; u++)
            R[v * m + u] = 0.;
    for (int i=1; i<p; i++) {
        BlockMatrixFreeTheta(&X[i * s * n], srht, s, P);
        cblas_dgemm(CblasColMajor, CblasTrans, CblasNoTrans, i * s, i * s, k, 1., S, k, S, k, 0., StS, i * s);
        cblas_dgemm(CblasColMajor, CblasTrans, CblasNoTrans, i * s, s, k, 1., S, k, P, k, 0., Rtmp, i * s);
        LAPACKE_dsysv(LAPACK_COL_MAJOR, 'U', i * s, s, StS, i * s, ipiv, Rtmp, i * s);
        for (int v=0; v<s; v++)
            for (int u=0; u<i*s; u++)
                R[i * s * m + v * m + u] = Rtmp[v * i * s + u];
        cblas_dcopy(n * s, &X[i * s * n], 1, &Q[i * s * n], 1);
        cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans, n, s, i * s, -1., Q, n, Rtmp, i * s, 1., &Q[i * s * n]);
        RGS(n, s, k, &Q[i * s * n], srht, &Q[i * s * n], Rtmp, &S[i * s * k]);
        for (int v=0; v<s; v++)
            for (int u=0; u<s; u++)
                R[i * s * m + v * m + i * s + u] = Rtmp[v * s + u];
        for (int v=0; v<s; v++)
            for (int u=(i+1)*s; u<m; u++)
                R[i * s * m + v * m + u] = 0.;
    }
    free(P);
    free(Rtmp);
    free(StS);
}
```


Stepping away from matrix computation

- ▶ SRHT refers to subsampled randomized Walsh-Hadamard transform.
- ▶ MatrixFreeTheta: $X \rightarrow \Theta X$, where Θ is a SRHT matrix given by:

$$\Theta := \sqrt{n/k} R H D$$

$R \in \mathbb{R}^{k \times n}$: Random restriction, i.e., each row is a row from I_n .

$H \in \mathbb{R}^{n \times n}$: Normalized Walsh-Hadamard transform matrix.

$D \in \mathbb{R}^{n \times n}$: Random sign flip, i.e., diagonal array with ± 1 components.

- ▶ We have $H = 1/\sqrt{n} H_n$, in which the non-normalized Walsh-Hadamard transform H_n is defined by the following recursion:

$$H_1 := \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}, \quad H_q := \begin{bmatrix} H_{q/2} & H_{q/2} \\ H_{q/2} & -H_{q/2} \end{bmatrix}, \quad q = 2, 4, \dots, n/2, n.$$

- ▶ The recurrence of the SRHT lends itself to **divide and conquer**, which yields a **non-vectorized fast algorithm**.

Algorithm for the fast Walsh-Hadamard transform (FWHT)

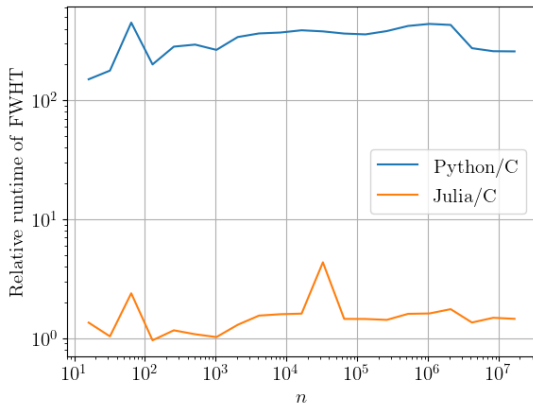
- Pseudocode of the FWHT assuming n is a power of 2:

Algorithm 2 FWHT : $z \mapsto H_n z$

```
1:  $h := 1$ 
2: while  $h < n$  do
3:   for  $i = 1, 1 + 2h, \dots, n - 2h, n$  do
4:     for  $j = i, \dots, i + h - 1$  do
5:        $x := z_j$ 
6:        $y := z_{j+h}$ 
7:        $z_j := x + y$ 
8:        $z_{j+h} := x - y$ 
9:    $h := 2h$ 
10: return  $z$ 
```

- If n is not a power of 2, zero-pad up to the smallest power of 2 greater than n .
- Then, define MatrixFreeTheta by making use of FWHT.

Runtime of the fast Walsh-Hadamard transform (FWHT)



C is approximately
1.5x faster than **Julia**
300x faster than **Python**

For **fast Python**:
code in **C** w/ **pybind11**
code in **Fortran** w/ **f2py**

Julia is fast
and **high level**

Other benchmarks

► See julialang.org/benchmarks/



Installing Julia

► Install juliaup:

- On Linux and MacOS:

```
$ curl -fsSL https://install.julialang.org | sh
```

- On Windows:

```
> winget install --name Julia --id 9NJNWW8PVKMN -e -s msstore
```

This will install the latest version of Julia.

► juliaup is also used to update Julia to the latest version:

```
$ juliaup update
```

► You can use juliaup to install arbitrary releases, e.g.:

```
$ juliaup add 1.9.3
```

► Start a new terminal and access the REPL as follows:

```
$ julia  
julia>
```

► To run a specific release, e.g.,

```
$ julia +1.9.3
```

► To see the installed versions: `$ juliaup status`

Package management

- ▶ The default project is defined by two files
 - `~/.julia/environments/v1.12/Project.toml`: contains names of packages.
 - `~/.julia/environments/v1.12/Manifest.toml`: contains version numbers and dependencies.
- ▶ To **clone your environment** on a new machine, only **copy** the **Project.toml** in the new default folder.
- ▶ In REPL, hit the key `]` to get in **Pkg mode**. You then get the following prompt:

```
(@1.12) pkg>
```

- ▶ To **initialize a new environment** in `~/MyEnvironment/`, **activate** the path and **add a package**:

```
(@1.12) pkg> activate ~/MyEnvironment/  
(@1.12) pkg> add NPZ
```

this will automatically create the `Project.toml` and `Manifest.toml` files in `~/MyEnvironment/`.

Package management, cont'd

- ▶ Load an **existing environment** using the **activate** command as follows:

```
(@1.12) pkg> activate ~/MyEnvironment/  
Activating project at '~/MyEnvironment'
```

- ▶ Check the status of an environment with the **st** command:

```
(@1.12) pkg> st  
Status '~/MyEnvironment/Project.toml'  
[15e1cf62] NPZ v0.4.2
```

- ▶ When loaded for the first time, use the command **instantiate** to **install** all the **packages** from **Project.toml**:

```
(@1.12) pkg> instantiate
```

- ▶ All these operations defined in the Pkg mode can be done with the Pkg package inside a script, e.g.:

```
using Pkg  
activate("~/MyEnvironment/")
```

Overview of available packages

► Scientific computing

- **LinearAlgebra.jl**: Basic linear algebra subroutines, multithreaded BLAS and LAPACK
- **SparseArrays.jl**: Support for sparse vectors and matrices
- **Distributed.jl**: Methods for distributed computing
- **DistributedArrays.jl**
- **MPI.jl**: Wrapper for the message passing interface
- **CUDA.jl**: Main entrypoint for programming NVIDIA GPUs
- **AlgebraicMultigrid.jl**: GPU-based implementation of AMG solvers and preconditioners
- **Metis.jl**: Wrapper for the Metis library
- **FFTW.jl**: Bindings to the FFTW library for fast Fourier transforms
- **SuiteSparse.jl**: Wrapper for the SuiteSparse library
- **Arpack.jl**: Wrapper for the Arpack library to solve large-scale eigenvalue problems
- **BenchmarkTools.jl**: Methods for performance tracking

Overview of available packages, cont'd

► Scientific computing (cont'd)

- **IterativeSolvers.jl**: Iterative algorithms to solve large linear systems
- **KrylovKit.jl**: Matrix-free Krylov-based algorithms for linear, singular value and eigenvalue problems
- **TriangleMesh.jl**: Generate and refine 2D unstructured triangular meshes
- **Gridap.jl**: Finite elements for partial differential equations in arbitrary dimensions
- **DifferentialEquations.jl**: Suite for numerically solving differential equations (including DAEs)

► Machine learning

- **Flux.jl**: Go-to library for neural networks and machine learning
- **Zygote.jl**: Automatic differentiation package
- **Knet.jl**: Deep learning framework developed at Koç University
- **TensorFlow.jl**: Wrapper for TensorFlow
- **ScikitLearn.jl**: Implementation of the scikit-learn API

Interoperability with Python

► Calling Python from Julia

- Set-up Python installation as follows using **PyCall.jl**:

```
julia> ENV["PYTHON"] = "/usr/bin/python3"
```

```
(@1.12) pkg> build PyCall
```

```
julia> using PyCall
```

- Import packages using **pyimport** from PyCall.jl:

```
julia> np = pyimport("numpy");
```

```
julia> pushfirst!(pyimport("sys")."path", "");
```

```
julia> GS = pyimport("GramSchmidt");
```

- Proceed seamlessly in Julia as in Python:

```
julia> x = np.random.rand(2^24);
```

```
julia> z = GS.fwht(x);
```

- Other packages:

- **PyPlot.jl**: Enables Matplotlib in Julia
- **NPZ.jl**: Enables saving and loading NumPy binary data files
- **Conda.jl**: Provides access to the Conda package manager

Interoperability with Python, cont'd

► Calling Julia from Python

- Install **PyJulia**:

```
$ pip install julia
```

- The **default environment** of Julia is then **available** from Python. For example, we can do

```
>>> from julia import NPZ
```

- The **global namespace** of Julia's interpreter can be accessed via the module **julia.Main**:

```
>>> from julia import Main
```

- You can set a variable's name in the **julia.Main** module to send data from Python to Julia:

```
>>> import numpy as np
```

```
>>> Main.x = np.random.rand(2**24)
```

- Use the **eval** function from **julia.Main** to run Julia code:

```
>>> Main.eval('push!(LOAD_PATH, ".")')      # add current
```

```
>>> Main.eval('using MyGramSchmidt: fwht')  # folder to path
```

```
>>> z = Main.eval('fwht(x)')
```

Interoperability with C

► Calling C libraries from Julia:

(see docs.julialang.org/en/v1/manual/calling-c-and-fortran-code/)

- The C code must be available as a **shared library**
- No additional overhead for calling from Julia compared to calling from C
- The function `ccall` is used to call a C function with the following arguments:
 - 1. A `(:function, "path/to/library")` pair
 - 2. The function's return type
 - 3. A tuple of input types corresponding to the function's signature
 - 4. Argument values to be passed to the function
- Example for `double *fwht(double *a, int n)` in `srht.so`:

```
julia> z_ptr = ccall((:fwht, "./srht.so"), Ptr{Cdouble},  
                    (Ptr{Cdouble}, Cint), rand(2^24),  
                    2^24)
```

```
julia> z = unsafe_wrap(Vector{Float64}, z_ptr, 2^24)
```

Interoperability with C, cont'd

► Calling Julia code from C:

(see docs.julialang.org/en/v1/manual/embedding/)

- A header file **julia.h** is made available in the Julia folder
- Example of C code (main.c) calling Julia code:

```
#include <julia.h>
```

```
JULIA_DEFINE_FAST_TLS
```

```
int main(int argc, char *argv[]) {
```

```
    jl_init();
```

```
    jl_eval_string("push!(LOAD_PATH, \".\")");
```

```
    jl_eval_string("using MyGramSchmidt: fwht");
```

```
    jl_array_t *z = (jl_array_t*)jl_eval_string("fwht(rand(2^24))")
```

```
    double *z_data = (double*)jl_array_data(z);
```

```
    jl_atexit_hook(0);
```

```
    return 0;}
```

- Compile as follows:

```
$ gcc -o main -fPIC
```

```
-I/home/venkovic/.julia/juliaup/julia-1.12.5+0.x64.linux.gnu/include/julia
```

```
-L/home/venkovic/.julia/juliaup/julia-1.12.5+0.x64.linux.gnu/lib
```

```
-Wl,-rpath,/home/venkovic/.julia/juliaup/julia-1.12.5+0.x64.linux.gnu/lib
```

```
main.c -ljulia
```

Shared memory multithreading

- ▶ The number of threads is set through an environment variable:

```
$ export JULIA_NUM_THREADS=12
```

- ▶ Multithreaded for loop:

```
julia> Z = Array{Float64,2}(undef, 1_024, Threads.nthreads())
julia> Threads.@threads for i in 1:Threads.nthreads()
    Z[:, i] = fwht(rand(1_024));
end
```

- ▶ Parallel task launching:

```
julia> a = Threads.@spawn fwht(rand(1_024));
julia> b = Threads.@spawn fwht(rand(1_024));
julia> z = fetch(a) .+ fetch(b)
```

- ▶ Multithreaded BLAS:

```
julia> using LinearAlgebra.BLAS
julia> BLAS.set_num_threads(Threads.nthreads())
julia> BLAS.dot(10_000_000, rand(10_000_000), 1,
    rand(10_000_000), 1);
```

Distributed computing

- ▶ Add aliases of your machines to `/etc/hosts`:

```
192.168.1.74 hector0 ... 192.168.1.23 hector3
```

- ▶ Set-up **password-less ssh** connection between machines

- ▶ Import the **Distributed.jl** package and add the machines:

```
using Distributed
machines = ["hector$i" for i in 0:3];
for machine in machines
    addprocs((["venkovic@$machine", Sys.CPU_THREADS]),
            tunnel=true)
end
```

- ▶ There are **processes** and **workers**. The **master process** is **not a worker**:

```
julia> println(procs(), workers())
[1,2,3,4,5,6,7,8,9,10,11,12] [2,3,4,5,6,7,8,9,10,11,12]
```

hector
4x Intel Core i7



Distributed computing, cont'd₁

- ▶ Load code on **all processes** making use of the **@everywhere** macro:

```
@everywhere push!(LOAD_PATH, ".")  
@everywhere using MyGramSchmidt: fwht
```

- ▶ Define a **shared array** as follows:

```
@everywhere using SharedArrays  
Z = SharedArray{Array{Float64,2}}(undef, 1_024, nworkers())
```

- ▶ Do a **distributed for loop** as follows. The loop is distributed over **workers**:

```
@distributed for p in 1:nworkers()  
    Z[:, p] = fwht(rand(1_024))  
end
```

- ▶ Use **pmap** as follows to divide the work among **workers**:

```
Z = pmap(i->fwht(rand(1_024)), 1:nworkers())
```

- ▶ Do a **reduction** through a **distributed for loop** as follows:

```
z = Array{Float64,1}(undef, 1_024)  
z .= @distributed (.) for _ in 1_nworkers()  
    fwht(rand(1_024))  
end
```


Distributed computing, cont'd₂

- **Dynamic mapreduce** routine for large parallel unbalanced working loads:
github.com/venkovic/julia-phd-krylov-spdes/blob/master/Utils/PIIUtils.jl

```
"""
```

```
function dynamic_mapreduce!(func::Function,  
                             redop::Function,  
                             coll::Array{Int,1},  
                             K::Array{Float64,2};  
                             verbose=true,  
                             Δt=2.)
```

```
Does parallel mapreduce of arrays with dynamic scheduling. This is an alternative to  
K .= @distributed (redop) for c in coll
```

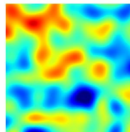
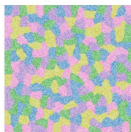
```
    func(c)  
end
```

```
which does static scheduling and tends to crash for time consuming and unbalanced work  
loads.
```

```
Another approach is given by reduce(redop, Distributed.pmap(func, K)), which requires to  
allocate enough memory to store Distributed.pmap(func, K). This becomes a problem when the  
number of workers and the dimensions of K are increased.
```

- Used for **parallel Karhunen-Loève decompositions on unstructured meshes**:

$n_d = 200$



Venkovic (2023)

Distributed computing, cont'd₃

- ▶ An alternative to reduction for loops is the **mapreduce** function:

```
z = mapreduce(x->fwht(rand(1_024)), .+, 1:nworkers())
```

- ▶ Launch a **task** on **any available worker**:

```
z = fetch(@spawn fwht(rand(1_024)))
```

- ▶ Launch a **task** on a **specific process**, say the 4th process:

```
z = fetch(@spawnat 4 fwht(rand(1_024)))
```

Performance tips

- ▶ Access arrays in memory order, i.e., **along columns**
- ▶ **Pre-allocate** returned variables

Instead of this:

```
for i in 1:10
    x = fwht(rand(1_024))
    println(x[1:3])
end
```

Use this:

```
x = Vector{Float64}(undef, 1_024)
for i in 1:10
    x[:] = fwht(rand(1_024))
    println(x[1:3])
end
```

- ▶ **Avoid changing the type** of a variable

Avoid this:

```
x = 1
for i in 1:10
    x *= rand()
end
```

Performance tips, cont'd

► Write **type-stable** functions

Instead of this: `pos(x) = x < 0 ? 0 : x`

Use this: `pos(x) = x < 0 ? zero(x) : x`

► Use **broadcast** operators for vectorized operations

Instead of this: `f(x::Vector{Float64}) = 3 * x.^2 + x`

Use this: `f(x::Vector{Float64}) = 3 .* x.^2 .+ x`

Useful macros

- ▶ Macros provide a mechanism to include generated code in the final body program. We've seen `@distributed`, but there are other examples:

- Use `@time` to time a command and get allocations info:

```
julia> @time fwht(rand(2^24));  
0.622758 seconds (4 allocations: 256.000 MiB, 1.34% gc time)
```

- Use `@elapsed` to store time elapsed during command execution:

```
julia> dt = @elapsed fwht(rand(2^24));  
julia> println("$dt seconds have passed.")  
0.624587781 seconds have passed.
```

- Use `@which` to identify the method invoked along with its signature and location in file:

```
julia> @which fwht(rand(1_024))  
fwht(a::Vector{Float64}) in MyGramSchmidt at ~/julia-gram-s
```

- Use `@code_llvm` to view the LLVM code used by the compiler:

```
julia> @code_llvm fwht(1_024)  
; @ ~/julia-gram-schmidt/GramSchmidt.jl:85 within 'fwht'  
define nonnull {}* @julia_fwht_819({}* nonnull align 16 der  
top:
```

Useful macros, cont'd

- Use `@code_native` to view the native assembly code generated by the compiler:

```
julia> @code_native fwht(rand(1_024));  
.text  
.file "fwht"  
.section      .rodata.cst8,"aM",@progbits,8  
.p2align      3          # -- Begin function julia_fwht_587  
.LCPI0_0  
...  
...
```

- Use `@code_warntype` to investigate type stability:

```
julia> @code_warntype fwht(rand(1_024));  
MethodInstance for fwht(::Vector{Float64})  
  from fwht(a::Vector{Float64}) in Main at REPL[16]:1  
Arguments  
  #self#::Core.Const(fwht)  
  a::Vector{Float64}  
Locals  
  ...
```


Ressources

- ▶ **Documentation:** docs.julialang.org
- ▶ **Discourse board:** discourse.julialang.org
Responsive community. Ideal for questions.
- ▶ **Slack:** julialang.slack.com
Good for package development, and questions.
- ▶ **YouTube:** www.youtube.com/c/TheJuliaLanguage
- ▶ Sengupta, Avik. *Julia High Performance: Optimizations, distributed computing, multithreading, and GPU programming with Julia 1.0 and beyond*. Packt Publishing Ltd, 2019.

My favorite



Used at Stanford



- ▶ Darve, Eric, and Mary Wootters. *Numerical Linear Algebra with Julia*. Vol. 172. SIAM, 2021.
- ▶ **JuliaCon 2026** @ Johannes Gutenberg University, Frankfurt, August 10-15, 2026.

References

- ▶ Balabanov, Oleg, and Laura Grigori. "Randomized block Gram-Schmidt process for solution of linear systems and eigenvalue problems." arXiv preprint arXiv:2111.14641 (2021).
- ▶ Balabanov, Oleg, and Laura Grigori. "Randomized Gram-Schmidt Process with Application to GMRES." SIAM Journal on Scientific Computing 44.3 (2022): A1450-A1474.

Homework

Homework

- ▶ Read chapter 10 - Julia Essentials in Darve and Wootters (2021)
- ▶ Watch [Julia Lightning Round on YouTube](#)