

Numerical Linear Algebra for Computational Science and Information Engineering

Lecture 09 Basic Iterative Methods for Eigenvalue Problems

Nicolas Venkovic
nicolas.venkovic@tum.de

Group of Computational Mathematics
School of Computation, Information and Technology
Technical University of Munich

Winter 2025-26



Outline I

- | | | |
|---|---|----|
| 1 | Methods for computing a single eigenvalue
Section 5.1 in Darve & Wootters (2021) | 4 |
| 2 | Basic QR iteration
Section 5.2 in Darve & Wootters (2021) | 15 |
| 3 | Other methods and implementations
Section 5.2 in Darve & Wootters (2021) | 27 |
| 4 | Homework problems | 30 |

Computing eigenvalues exactly is impossible

- ▶ Computing eigenvalues and eigenvectors is a very difficult task.

There is no direct method for computing eigenvalues of matrices of size five or higher in general.

That is, there is no algorithm that can compute eigenvalues exactly assuming exact arithmetic.
- ▶ Moreover, it can be proved that a method that computes eigenvalues exactly cannot exist for general matrices of size five or higher.

The reason for this is the Abel-Ruffini theorem, which states that no direct method exists to find exact zeros of a polynomial of degree five or higher.

That is the case because computing the roots of any polynomial is equivalent to finding the eigenvalues of a matrix.

Thus, since there is no method for finding zeros of a polynomial, then there cannot exist an exact method for finding eigenvalues of a general matrix.

Computing eigenvalues exactly is impossible, cont'd

You saw one side of the equivalence between solving for eigenvalues of a general matrix and solving for the zeros of a polynomial in your Linear Algebra class.

To see the other direction, consider a generic polynomial given by

$$p(x) = x^n + a_{n-1}x^{n-1} + \cdots + a_1x + a_0.$$

Then, there is a matrix

$$A = \begin{pmatrix} 0 & 1 & & & & \\ & 0 & 1 & & & \\ & & 0 & & & \\ & \vdots & & & & \\ & & & & 0 & 1 \\ -a_0 & -a_1 & -a_2 & \cdots & -a_{n-2} & -a_{n-1} \end{pmatrix}$$

such that, if we pick $u = [1 z z^2 \dots z^{n-1}]^T$ where z is a root of $p(x)$, then we have $Au = zu$ so that (z, u) is an eigenpair of A .

Consequently, all roots of $p(x)$ are eigenvalues of A .

Convention

- ▶ Let us denote $A = X\Lambda X^{-1}$ an eigendecomposition of A .

In this lecture, all the algorithms will normalize vectors, i.e., replace x by $x/\|x\|_2$ during the iterative process.

Therefore, when discussing convergence, we will assume the columns of X have norm 1.

This is done without loss of generality, since $A = X\Lambda X^{-1}$ remains valid irrespective of the magnitude of the columns of X .

Moreover, in many places, results will be stated "up to a sign" or "up to a unit complex factor", because even with normed columns, the matrix X of an eigendecomposition is not unique.

Methods for computing a single eigenvalue

Section 5.1 in Darve & Wootters (2021)

Taking powers of A

- ▶ Suppose that A is a square diagonalizable matrix.

Then A has an eigenvalue decomposition $A = X\Lambda Y^H$ where the columns x_i of X are right eigenvectors of A , and the columns y_i of $Y := X^{-H}$ are left eigenvectors of A :

$$A = \begin{matrix} \text{green square} \\ x_i \downarrow \\ X \quad \Lambda \quad X^{-1} \end{matrix} = \begin{matrix} \text{blue vertical bar} \\ \text{blue rectangle} \end{matrix} \quad \begin{matrix} \text{white rectangle with red diagonal} \\ \text{blue rectangle} \end{matrix} \quad \begin{matrix} \leftarrow y_i^H \end{matrix}$$

Darve, E., & Wootters, M. (2021). Numerical linear algebra with Julia. Society for Industrial and Applied Mathematics.

One thing about the eigendecomposition is that powers of A are such that

$$A^k = X\Lambda^k Y^H = \sum_i \lambda_i^k x_i y_i^H.$$

Notice that, even if A is real, it can have complex eigenvalues and vectors.

Note also that left and right eigenvectors of A coincide if A is normal.

Taking powers of A , cont'd

- Let us assume the eigenvalues of A are ordered such that

$$|\lambda_1| > |\lambda_2| \geq \cdots \geq |\lambda_n|$$

where, in particular, the largest eigenvalue has magnitude strictly greater than the second one.

Then, even for moderate values of the power k , we expect λ_1^k to dominate in A^k , i.e., $|\lambda_1^k| \gg |\lambda_2^k| \geq \cdots \geq |\lambda_n^k|$ so that

$$A^k = \lambda_1^k x_1 y_1^H + \cdots + \lambda_n^k x_n y_n^H \approx \lambda_1^k x_1 y_1^H.$$

- Let's multiply A^k by a random vector z , such that $y_1^H z$ is not too small, then

$$A^k z \approx \lambda_1^k x_1 y_1^H z = \lambda_1^k (y_1^H z) x_1$$

so that $A^k z / \|A^k z\|_2$ gives a **good approximation of x_1** .

Power iteration

- The power iteration is based on this idea of taking powers of A to approximate the largest eigenpair. The algorithm is as follows:

1. Sample a random vector $q^{(0)} \in \mathbb{C}^n$

2. $q^{(0)} := q^{(0)} / \|q^{(0)}\|_2$

3. For $k = 0, 1, 2 \dots$

4. $z^{(k)} := Aq^{(k)}$

5. $\lambda^{(k+1)} = z^{(k)H} q^{(k)}$

6. $q^{(k+1)} := z^{(k)} / \|z^{(k)}\|_2$

where $(\lambda^{(k)}, q^{(k)})$ is an iterate approximating the largest eigenpair of A .

At the k -th step, the approximate eigenvector is

$$q^{(k)} = A^k q^{(0)} / \|A^k q^{(0)}\|_2,$$

and the corresponding approximate eigenvalue is $\lambda^{(k)} = q^{(k)H} A q^{(k)}$.

Note that, even though $q^{(k)}$ is formed with A^k , the matrix power A^k is not explicitly computed.

Instead, we just perform repeated matrix-vector products.

Convergence of power iteration

- ▶ Let us assume again that the eigenpairs $(\lambda_1, x_1), \dots, (\lambda_n, x_n)$ of A are ordered such that $|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_n|$.
- ▶ The starting vector $q^{(0)}$ can be expressed in the basis formed by the eigenvectors of A , i.e.,

$$q^{(0)} = \alpha_1 x_1 + \dots + \alpha_n x_n.$$

For the method to work, we need to assume $\alpha_1 \neq 0$, that is, $q^{(0)}$ is not orthogonal to x_1 .

- ▶ Then, we have

$$A^k q^{(0)} = \sum_{i=1}^n \alpha_i A^k x_i = \sum_{i=1}^n \alpha_i \lambda_i^k x_i$$

which can be factorized as follows:

$$\begin{aligned} A^k q^{(0)} &= \alpha_1 \lambda_1^k x_1 + \alpha_2 \lambda_2^k x_2 + \dots + \alpha_n \lambda_n^k x_n \\ &\quad \alpha_1 \lambda_1^k \left(x_1 + \frac{\alpha_2}{\alpha_1} \left(\frac{\lambda_2}{\lambda_1} \right)^k x_2 + \dots + \frac{\alpha_n}{\alpha_1} \left(\frac{\lambda_n}{\lambda_1} \right)^k x_n \right) \end{aligned}$$

Convergence of power iteration, cont'd

From that expression, we have

$$\|A^k q^{(0)}\|_2 = |\alpha_1 \lambda_1^k| (1 + \mathcal{O}(|\lambda_2/\lambda_1|)) \text{ and}$$

$$\|(\alpha_1 \lambda_1^k)^{-1} A^k q^{(0)} - x_1\|_2 = \mathcal{O}\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right)$$

which, along with the fact that $\|A^k q^{(0)}\|_2 \approx |\alpha_1 \lambda_1^k|$ implies that our estimate $q^{(k)} = A^k q^{(0)} / \|A^k q^{(0)}\|$ approaches x_1 with an error $\mathcal{O}(|\lambda_2/\lambda_1|^k)$.

In summary, we have

$$\|q^{(k)} - x_1\|_2 = \mathcal{O}\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right) \quad \text{and} \quad |\lambda^{(k)} - \lambda_1| = \mathcal{O}\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right).$$

- ▶ Although it is a good starting point, this version of power iteration is limited as it cannot find approximates of any eigenvalue except the largest one. It also cannot leverage given approximations of λ_i .

Inverse iteration

- ▶ Assume we are equipped with an approximation μ of the eigenvalue λ_i of A .
- ▶ An **inverse iteration** uses μ to form an arbitrarily good approximation of λ_i .
- ▶ If μ is a good approximation of λ_i , then

The shifted matrix $A - \mu I_n$ has a small eigenvalue $\lambda_i - \mu$.

The shift-and-invert matrix $(A - \mu I_n)^{-1}$ has a large eigenvalue $1/(\lambda_i - \mu)$.

So, a power iteration applied to $(A - \mu I_n)^{-1}$ should allow us to calculate x_i very quickly, since $1/(\lambda_i - \mu)$ is now the largest eigenvalue, with the corresponding eigenvector x_i .

The algorithm of inverse iteration is as follows:

1. Sample a random vector $q^{(0)} \in \mathbb{C}^n$
2. $q^{(0)} := q^{(0)} / \|q^{(0)}\|_2$
3. For $k = 0, 1, 2 \dots$
4. Solve for $z^{(k)}$ s.t. $(A - \mu I_n)z^{(k)} = q^{(k)}$ // $z^{(k)} := (A - \mu I_n)^{-1}q^{(k)}$
5. $q^{(k+1)} := z^{(k)} / \|z^{(k)}\|_2$
6. $\lambda^{(k+1)} = q^{(k+1)H} A q^{(k+1)}$

Convergence of inverse iteration

- ▶ Similarly to power iteration, we can characterize the convergence of inverse iterations by

$$|\lambda^{(k)} - \lambda_i| = \mathcal{O} \left(\left| \frac{\lambda_i - \mu}{\lambda_j - \mu} \right|^k \right)$$

where λ_i and λ_j are the closest and second closest eigenvalues of A to μ , respectively.

If $|\lambda_i - \mu| \ll |\lambda_j - \mu|$, then the convergence is fast.

Rayleigh quotient iteration

- ▶ As inverse iterations progress, the iterate $\lambda^{(k)}$ becomes a better approximation of the eigenvalue λ_i than μ .
One could use this fact to redefine the shift μ and get faster convergence.
- ▶ Let us assume the matrix A is real and symmetric so that its eigenvalues and eigenvectors are real, and the eigenvectors are orthogonal.
- ▶ The idea to **update the shift μ during the iteration** is deployed in an algorithm called **Rayleigh quotient iteration**.

Let us consider the **Rayleigh quotient** given by $r(x) = \frac{x^T A x}{x^T x}$ for $x \neq 0$.

The Rayleigh quotient is used to approximate an eigenvalue.

Indeed, note that if x is an eigenvector of A , i.e., $Ax = \lambda x$, then $r(x) = \lambda$ is the corresponding eigenvalue.

Rayleigh quotient iteration, cont'd

- The algorithm for Rayleigh quotient iterations is as follows:

1. Sample a random vector $q^{(0)} \in \mathbb{C}^n$
2. $q^{(0)} := q^{(0)} / \|q^{(0)}\|_2$
3. $\lambda^{(0)} := \mu$
4. For $k = 0, 1, 2 \dots$
5. Solve for $z^{(k)}$ such that $(A - \lambda^{(k)} I_n) z^{(k)} = q^{(k)}$
6. $q^{(k+1)} := z^{(k)} / \|z^{(k)}\|_2$
7. $\lambda^{(k+1)} = q^{(k+1)H} A q^{(k+1)}$

- Rayleigh quotient iterations converge faster than inverse iterations.

Convergence of Rayleigh quotient iterations

- ▶ Note first that the gradient of the Rayleigh quotient r for a symmetric A is given by $\nabla r(x) = \frac{2}{x^T x} (Ax - r(x)x)$ so that $r(x_i) = \lambda_i$ implies $\nabla r(x_i) = 0$. More often than not, the zeros of ∇r are saddle points, as the Rayleigh quotient is only minimized (resp. maximized) at the smallest eigenpair (resp. largest eigen-pair).

In particular, we remember the Courant-Fischer theorem from lecture 1 which states

$$\lambda_{\min} = \min_{x \neq 0} \frac{x^T Ax}{x^T x} \quad \text{and} \quad \lambda_{\max} = \max_{x \neq 0} \frac{x^T Ax}{x^T x}.$$

- ▶ Then, suppose that y is close to an eigenvector x_i , by Taylor expansion around x_i , we have

$$r(y) \approx r(x_i) + \nabla r(x_i)^T(y - x_i) + (y - x_i)^T H(x_i)(y - x_i)$$

$r(x_i) = \lambda_i$

This is zero since
 $\nabla r(x_i) = 0$

Here, H is the Hessian matrix
We have:

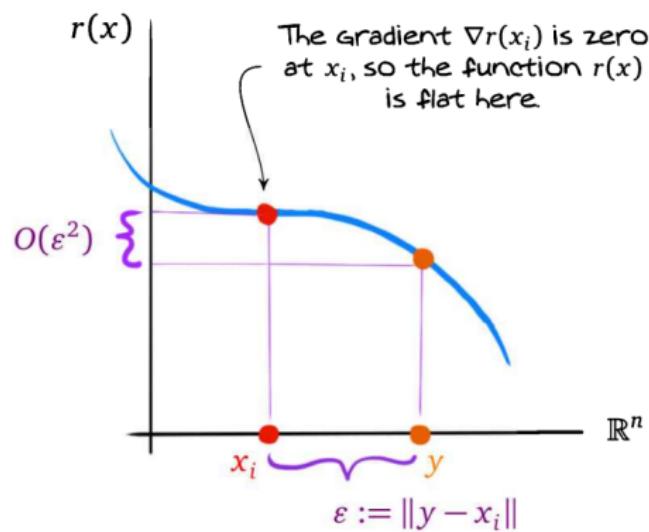
$$(y - x_i)^T H(x_i)(y - x_i) \leq \|H(x_i)\| \|y - x_i\|^2 = O(\|y - x_i\|^2)$$

Convergence of Rayleigh quotient iterations, cont'd

Consequently, the first order term disappears, leaving us with

$$r(y) = \lambda_i + \mathcal{O}(\|y - x_i\|_2^2)$$

and the behavior of the Rayleigh quotient near an eigenvector x_i is as follows:



Basic QR iteration

Section 5.2 in Darve & Wootters (2021)

Basic QR iteration

- ▶ The PageRank algorithm is a variant of power iteration aimed at finding the largest eigenvector of a modified adjacency matrix of a web graph. However, in general, iterative methods for computing a single eigenpair have limited applicability.
- ▶ Unlike those previously covered iterative methods for eigenvalue solving, QR iterations aim at **finding all the eigenvalues** of a matrix.
- ▶ The QR iteration was elected **one of the 10 best algorithms of the 20th century** by Dongarra and Sullivan (2000).
- ▶ The QR iteration is the **state of the art eigensolver for small dense eigenvalue problems**. It is implemented in LAPACK, and it serves as a **building block of larger**, possibly sparse **iterative eigensolvers**.
- ▶ An important assumption of this Section is that A is **diagonalizable with separate eigenvalues**, i.e., such that $|\lambda_1| > |\lambda_2| > \dots > |\lambda_n|$.
Because A is real with separate eigenvalues, we have that eigen- and Schur decompositions of A are real.

Dongarra, J., & Sullivan, F. (2000). Guest editor's introduction: The top 10 algorithms. Computing in Science & Engineering 2, 22–23.

Orthogonal iteration for $r = 2$

- ▶ Orthogonal iterations allow us to recover more than one eigenvalue at once.
- ▶ For starters, consider that only $r = 2$ eigenvalues are needed.
Then, the pseudocode of orthogonal iterations is as follows:

1. Sample two random vectors $q_1, q_2 \in \mathbb{R}^n$
2. While not converged :
3. $q_1 := Aq_1; q_2 := Aq_2$
4. Project q_2 onto the space orthogonal to q_1 // $q_2 := \left(I_n - \frac{q_1 q_1^T}{q_1^T q_1} \right) q_2$
5. $q_1 := q_1 / \|q_1\|_2, q_2 := q_2 / \|q_2\|_2$
6. Return $q_1^T A q_1$ and $q_2^T A q_2$

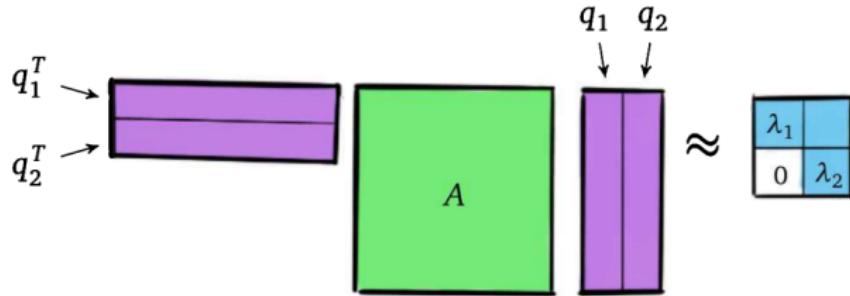
Disregarding the vector q_2 , the vector q_1 undergoes a standard power iteration so that, at the k -th step, we have

$$q_1^{(k)} = \frac{A^k q_1^{(0)}}{\|A^k q_1^{(0)}\|_2}$$

which converges towards x_1 .

Orthogonal iteration for $r = 2$, cont'd₁

- If we assume that $q_1 \approx x_1$ has already converged, then the update step for q_2 is of the form
$$q_2^{(k)} \approx (I_n - x_1 x_1^T) A q_2^{(k-1)}$$
 where $I_n - x_1 x_1^T$ is the orthogonal projector onto $\text{span}\{x_1\}^\perp$. Thus, q_2 is undergoing a power iteration with the matrix $(I_n - x_1 x_1^T) A$. It can be shown that the largest eigenvalue of this matrix is λ_2 with an eigenvector along $(I_n - x_1 x_1^T)x_2$ towards which q_2 converges.
- Note that, if x_1 and x_2 are not orthogonal, then $(I_n - x_1 x_1^T)x_2$ is not aligned with x_2 . However, we do have $\text{span}\{q_1, q_2\} = \text{span}\{x_1, x_2\}$. Then, we claim that $Q^T A Q$ where $Q = [q_1 \ q_2]$ converges to an upper-triangular matrix with λ_1 and λ_2 on the diagonal :



Orthogonal iteration for $r = 2$, cont'd₂

- First, the upper-triangularity is explained as follows:

$$q_2^T A q_1 \approx q_2^T A x_1 = \lambda_1 q_2^T x_1 \approx \lambda_1 q_2^T q_1 = 0$$

so that the lower-left entry converges to zero.

- To see that λ_1 and λ_2 lie on the diagonal, it suffices to show that they are eigenvalues of $Q^T A Q$, as $Q^T A Q$ is triangular.

For this, since $\text{span}\{q_1, q_2\} = \text{span}\{x_1, x_2\}$ after convergence, then there is $v_i \in \mathbb{R}^2$ such that $Qv_i \approx x_i$ and we have

$$Q^T A Q v_i \approx Q^T A x_i = \lambda_i Q^T x_i \approx \lambda_i Q^T Q v_i = \lambda_i v_i$$

so that v_i is an eigenvector of $Q^T A Q$ with eigenvalue λ_i for $i = 1, 2$.

- Since Q is orthogonal and $Q^T A Q$ is upper triangular with the same eigenvalues as A , it seems that $Q(Q^T A Q)Q^T$ is a Schur decomposition of A .

Orthogonal iteration for general r

- When an arbitrary number r of eigenvalues is sought, the approximate eigenvectors are orthogonalized by performing a QR factorization, leading to the following pseudocode:

1. Sample a random matrix $Q_0 \in \mathbb{R}^{n \times r}$
2. $k := 0$
3. While not converged :
4. $Y_{k+1} := AQ_k$
5. Compute QR factorization $Q_{k+1}R_{k+1} = Y_{k+1}$
6. $k := k + 1$
7. Return $\text{diag}(Q_k^T AQ_k)$

Similarly as with $r = 2$, this method converges to an upper-triangular matrix $Q_k^T AQ_k$ with eigenvalues $\lambda_1, \dots, \lambda_r$.

Once the algorithm has converged, the approximate eigenvalues can be read from the diagonal of the Schur form $Q_k^T AQ_k$.

Convergence of orthogonal iteration for general r

- ▶ If A is symmetric, then the eigenvectors x_1, \dots, x_r are orthogonal, and the i -th column of Q_k , which we denote by $q_i^{(k)}$, converges to $\pm x_i$.
For general matrices, things are different.
- ▶ Let us denote the matrices $Q^x \in \mathbb{R}^{n \times r}$ and $R^x \in \mathbb{R}^{r \times r}$ such that

$$[x_1 \dots x_r] = Q^x R^x.$$

We see that the iterate Q_k converges to Q^x :

Since $q_1^{(k)}$ undergoes a normal power iteration, it converges to $x_1 = q_1^x$.

For $q_2^{(k)}$, the QR decomposition ensures $\text{span}\{x_1, x_2\} = \text{span}\{q_1^x, q_2^x\}$ and we have

$$\text{span}\{q_1^{(k)}, q_2^{(k)}\} \approx \text{span}\{x_1, x_2\} = \text{span}\{q_1^x, q_2^x\}$$

Thus $q_2^{(k)}$ converges to something in the space $\text{span}\{q_1^x, q_2^x\}$, and it also has to be orthogonal to $q_1^{(k)} \approx q_1^x$. Therefore $q_2^{(k)}$ has to converge to $\pm q_2^x$. Similarly, $q_i^{(k)}$ converges to $\pm q_i^x$. Overall, we have that Q_k converges to Q^x .

Convergence to the Schur decomposition

- ▶ Now that we know that Q_k converges to Q^x , we can analyze the matrix $Q_k^T A Q_k$, which converges to $Q^{xT} A Q^x$ up to some columnwise sign changes.
- ▶ Since $AX = X\Lambda$ where $X = [x_1, \dots, x_r]$ and $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_r)$, the definitions of Q^x and R^x imply that

$$AX = X\Lambda$$

$$AQ^x R^x = Q^x R^x \Lambda$$

$$Q^{xT} A Q^x R^x (R^x)^{-1} = Q^{xT} Q^x R^x \Lambda (R^x)^{-1}$$

$$Q^{xT} A Q^x = R^x \Lambda (R^x)^{-1}.$$

Since R^x is upper triangular and Λ is diagonal, we have that $R^x \Lambda (R^x)^{-1}$ is upper triangular.

More particularly, we also have that $Q^{xT} A Q^x$ is upper triangular with the eigenvalues $\lambda_1, \dots, \lambda_r$ on the diagonal.

Then, the matrix $Q_k^T A Q_k$ converges to $Q^{xT} A Q^x$, which is upper triangular and has the top r eigenvalues of A on the diagonal.

Convergence of orthogonal iteration

- ▶ The convergence analysis being sequential, i.e., we assumed $q_1^{(k)} \approx q_1^x$, then showed that $q_2^{(k)}$ converges to q_2^x , and so on; may lead to think that the convergence of orthogonal iteration is slow. I.e., we first have to wait that $q_1^{(k)}$ converges, then $q_2^{(k)}$, and so on.
But, in fact, what actually happens is that all of the $q_i^{(k)}$ converge simultaneously.
- ▶ It can be shown that the convergence of the iterate Q_k to Q^x depends, similarly as before, on the separation between λ_r and λ_{r+1} . In particular, we have

$$\|Q_k Q_k^T - Q^x Q^{xT}\|_2 = \mathcal{O} \left(\left| \frac{\lambda_{r+1}}{\lambda_r} \right|^k \right).$$

That is, the smaller $|\lambda_{r+1}/\lambda_r|$, the faster the convergence of Q_k to Q^x .

QR iteration

- ▶ QR iterations are a re-framing of orthogonal iterations with $r = n$.
QR iterations yield the full Schur decomposition $T = Q^T A Q$ of A where T is an n -by- n upper triangular matrix with the eigenvalues of A on the diagonal, and Q is a n -by- n orthogonal matrix of a QR decomposition of the eigenvectors X of A .
- ▶ The iterate of QR iteration is denoted by Q_k with a corresponding matrix $T_k := Q_k^T A Q_k$.
- ▶ The formulation of QR iterations is more commonly expressed as a recurrence from $T_k = Q_k^T A Q_k$ to $T_{k+1} = Q_{k+1}^T A Q_{k+1}$.

From the definition of orthogonal iterations, we have

$$Q_{k+1} R_{k+1} = A Q_k \text{ so that } T_k = Q_k^T A Q_k = Q_k^T Q_{k+1} R_{k+1}$$

and, since $r = n$, we have $Q_k Q_k^T = I_n$ and

$$R_{k+1} Q_k^T = Q_{k+1}^T A \text{ so that } T_{k+1} = Q_{k+1}^T A Q_{k+1} = R_{k+1} Q_k^T Q_{k+1}$$

QR iteration, cont'd₁

Then, as we let $U_{k+1} := Q_k^T Q_{k+1}$, we have

$$T_k = U_{k+1} R_{k+1}$$

$$T_{k+1} = R_{k+1} U_{k+1}$$

where R_{k+1} is upper triangular, and U_{k+1} is orthogonal.

Note that $U_{k+1} R_{k+1}$ is a QR decomposition of $T_k = Q_k^T A Q_k$.

This yields the following pseudocode to compute the eigenvalues of A :

1. $T_0 := A$
2. $k := 0$
3. While not converged :
4. Compute QR factorization $U_{k+1} R_{k+1} = T_k$
5. $T_{k+1} := R_{k+1} U_{k+1}$
6. $k := k + 1$
7. Return $\text{diag}(T_k)$

Notice that A is only needed at the start of the algorithm, after what we only repeatedly compute QR decompositions and switch the factors.

QR iteration, cont'd₂

- In this algorithm, the matrix $U_{k+1} = Q_k^T Q_{k+1}$ represents an orthogonal correction.

Since upon convergence $U_k \rightarrow I_n$, the determinant of U_k is 1 for large k , and we can interpret U_k as a small rotation on the orthogonal vectors in Q_k . In particular, we have:

$$U_1 \dots U_{k+1} = Q_0^T Q_1 Q_1^T Q_2 \dots Q_k^T Q_{k+1} = Q_0 Q_{k+1} = Q_{k+1}$$

because we chose $Q_0 = I_n$.

As the algorithm converges, Q_k and Q_{k+1} become very close.

- In the symmetric case, $T_k = Q_k^T A Q_k$ is symmetric, but since it also converges to an upper symmetric matrix, it actually converges to a diagonal form, in which case the Schur decomposition is actually an eigendecomposition.

QR iteration, cont'd₃

- ▶ The QR iteration presented so far has drawbacks:
 - A QR factorization at cost $\mathcal{O}(n^3)$ is computed at each iteration.
 - The convergence depends heavily on the distribution of the eigenvalues, and it may never converge if two eigenvalues have the same magnitude.
- ▶ Improvements of the QR iteration method can be introduced to improve the robustness and efficiency:
 - The transformation of A into an upper Hessenberg form allows to decrease the cost of the QR factorizations.
 - A shifted version of the QR iteration can improve convergence, even when the eigenvalues are not well-separated, making the method robust to cases of eigenvalues with equal magnitudes.

Other methods and implementations

Section 5.2 in Darve & Wootters (2021)

Divide-and-conquer method

- ▶ A symmetric matrix can efficiently be transformed into a tridiagonal form using an orthogonal transformation

$$Q^T A Q = T.$$

Then, the eigendecomposition of A can be obtained from that of T .

- ▶ The divide-and-conquer method splits the tridiagonal matrix into two tridiagonal blocks plus a rank-1 perturbation:

$$T = \begin{bmatrix} T_1 & \\ & T_2 \end{bmatrix} + \rho u u^T.$$

- ▶ The method proceeds as follows:
 - ① Calculate the eigendecompositions of T_1 and T_2 .
 - ② The rank-1 perturbation allows to compute the eigenvalues of T given the eigendecompositions of T_1 and T_2 .

Method of bisection

- ▶ The method of bisection also considers a tridiagonal form $Q^T A Q = T$.
- ▶ The eigenvalues of T are the roots of $p_n(\lambda) = \det(T - \lambda I_n)$.

Finding these roots is generally a complex problem, but it can be simplified if we consider only the leading r -by- r block T_r of T and the corresponding characteristic polynomial

$$p_r(\lambda) = \det(T_r - \lambda I_r).$$

- ▶ As T is tridiagonal, it is possible to find a simple relation between p_r , p_{r-1} and p_{r-2} .

Using this sequence of polynomials, the method of bisection is able to efficiently calculate the roots of p_n .

Existing implementations

- ▶ QR iteration:
 - Available for general matrices.
 - Implementation sometimes requires tridiagonalization.
 - Fastest to compute the eigendecomposition of small matrices ($n \leq 25$).
 - Algorithm behind the Matlab, NumPy and Julia functions.
 - Available in LAPACK as `ssyev` for dense symmetric matrices.
 - Available in LAPACK as `sstev` for symmetric tridiagonal matrices.
- ▶ Divide-and-conquer method:
 - Available for symmetric matrices.
 - Implementation requires tridiagonalization.
 - Fastest to compute the eigendecomposition of medium size tridiagonal matrices, i.e., for $n > 25$.
 - Available in LAPACK as `sstevd` for symmetric tridiagonal matrices, `sstevd` defaults to QR iteration for smaller matrices.
- ▶ Method of bisection:
 - Available in LAPACK as `ssyevx` for dense symmetric matrices.

Homework problems

Homework problem

Turn in **your own** solution to **Pb. 20**:

Pb. 19 Show that the gradient of the Rayleigh quotient of a symmetric matrix A given by

$$r(x) = \frac{x^T A x}{x^T x} \text{ for } x \neq 0$$

is $\nabla r(x) = \frac{2}{x^T x} (Ax - r(x)x)$.

Pb. 20 Let (λ, x) be a right eigen-pair of A , (μ, y) be a left eigen-pair of A and μ be distinct from λ . Show that x and y are orthogonal.