

Numerical Linear Algebra for Computational Science and Information Engineering

Lecture 06 Introduction to Direct Methods for Sparse Linear Systems

Nicolas Venkovic
nicolas.venkovic@tum.de

Group of Computational Mathematics
School of Computation, Information and Technology
Technical University of Munich

Winter 2025-26



Outline I

- | | | |
|---|---|----|
| 1 | Solving sparse triangular linear systems
Section 9.3 in Darve & Wotters (2021) | 1 |
| 2 | Cholesky factorization
Section 9.4 in Darve & Wotters (2021) | 5 |
| 3 | Nested dissection
Section 9.5 in Darve & Wotters (2021) | 20 |
| 4 | Homework problems | 26 |

Solving sparse triangular linear systems

Section 9.3 in Darve & Wotters (2021)

When L is sparse and b is dense

- ▶ We want to solve $Lx = b$ where L is a **sparse lower-triangular** matrix with **non-zero diagonal** entries.
- ▶ Remember how to proceed when L is dense:

1. $x_1 = b_1/l_{11}$

2. $x_2 = (b_2 - l_{21}x_1)/l_{22}$

\vdots

$i.$ $x_i = \left(b_i - \sum_{j=1}^{i-1} l_{ij}x_j\right)/l_{ii}$

\vdots

$n.$ $x_n = \left(b_n - \sum_{j=1}^{n-1} l_{nj}x_j\right)/l_{nn}$

- ▶ When L is sparse, we simply need to **skip** the **zero components** l_{ij} in **each summand**.
- ▶ We will see in practice session that this can easily be implemented.
- ▶ The **final form of the implementation depends on** the sparse matrix **data structure** used to store L .

When both L and b are sparse

► When L and b are sparse, then the solution x may be sparse.

► Ideally, we would like to solve for x as follows:

1. for $i = 1, \dots, n$:
2. if $x_i \neq 0$:
3. $x_i \leftarrow b_i / \ell_{ii}$
4. for $j = 1, \dots, i - 1$:
5. if $\ell_{ij} \neq 0$:
5. $x_i \leftarrow x_i - \ell_{ij} x_j / \ell_{ii}$

► But iterating over the non-zero components of x requires to know the structure of x .

► For any non-zero x_i , we have either or both

(a) $b_i \neq 0$

(b) there is some $j < i$ such that $\ell_{ij} \neq 0$ and $x_j \neq 0$.

When both L and b are sparse, cont'd₁

- Let $G = (V, E)$ be the graph associated with L , then we denote by $X \subset V$ the minimal set of vertices so that either or both (a) and (b) hold.


That is, $X \subset V$ is the minimal set such that:

$$b_i \neq 0 \implies i \in X \text{ and } \ell_{ij} \neq 0 \text{ and } j \in X \implies i \in X.$$

Definition (Reachability & Reach)

- A vertex $i \in V$ in a directed graph $G = (V, E)$ is **reachable from a vertex** $j \in V$, if there is a directed path from j to i in G . That is, if there is a sequence of edges $(j, i_1), (i_1, i_2), \dots, (i_{k-1}, i_k), (i_k, i)$ where all the edges are in E .
- The set of vertices $i \in V$ reachable from a vertex $j \in V$ is the **reach** of j .

- If $j \in X$, then every vertex in the reach of j , is also in X .



When both L and b are sparse, cont'd₂

- ▶ Then, if we let $B \subset V$ be the set of vertices $i \in V$ such that $b_i \neq 0$, then X is the **set of vertices reachable from B** .
- ▶ Consequently, the set X can be found by operating on the graph $G = (V, E)$ associated with L .

Namely, the set X can be found using a **depth-first traversal (DFS, i.e., for depth-first search)** from every vertex in B .

- ▶ Depth-first traversal starts from some node j , and explores as far as possible along each branch in the graph before backtracking.
- ▶ We will see an implementation of depth-first traversal in the practice session.
- ▶ Procedure to solve $Lx = b$ where both L and b are sparse is as follows:

Linear solve of $Lx = b$ where both L and b are sparse

1. Define the set B from the sparsity pattern of b .
2. Find the set X of non-zero x components using DFS on B .
3. Run modified version of the algorithm to solve $Lx = b$ with a sparse L , but compute x_i only if $i \in X$.

Cholesky factorization

Section 9.4 in Darve & Wotters (2021)

Up-looking Cholesky algorithm

- ▶ Now that we know how to solve sparse triangular systems, we can use this to obtain a sparse Cholesky factorization.
- ▶ In particular, the **up-looking Cholesky** algorithm performs a Cholesky factorization by doing a **series of sparse triangular solves**.
- ▶ Proceeding by construction, assume the $(n - 1)$ -dimensional leading block L' of the Cholesky factor L of A is already known, leading the following structure of the $LL^T = A$ factorization:

$$\begin{bmatrix} L' & 0 \\ x^T & w \end{bmatrix} \begin{bmatrix} (L')^T & x \\ 0 & w \end{bmatrix} = \begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$$

First, we have $\begin{bmatrix} L' & 0_{(n-1) \times 1} \end{bmatrix} \begin{bmatrix} x \\ w \end{bmatrix} = b$ which simplifies to $L'x = b$.

Second, we have $\begin{bmatrix} x^T & w \end{bmatrix} \begin{bmatrix} x \\ w \end{bmatrix} = a$ and $w > 0$ so that $w = \sqrt{a - x^T x}$.

Up-looking Cholesky algorithm, cont'd₁

► This leads to the following algorithm:

Up-looking Cholesky algorithm

Given a sparse SPD matrix $A \in \mathbb{R}^{n \times n}$, initialize $L' := \sqrt{a_{11}}$.

For $k = 2, \dots, n$:


- Let the leading k -by- k block of A be written as $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ where A' is the $(k-1)$ -dimensional leading block, b is $(k-1)$ -by-1 and a is a scalar.
- Solve for $x \in \mathbb{R}^{k-1}$ such that $L'x = b$ where L' and b are sparse.
- Compute $w := \sqrt{a - x^T x}$, and update

$$L' := \begin{bmatrix} L' & 0 \\ x^T & w \end{bmatrix}$$

Return $L := L'$


Consequently, the sparse Cholesky factor L of the sparse matrix A is formed by performing $n - 1$ sparse triangular solves of sizes $1, \dots, n - 1$.

Up-looking Cholesky algorithm, cont'd₂



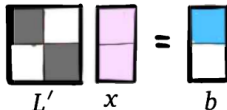
- Consider a matrix A with the following non-zero pattern:

- Since the 2-by-2 leading block A' of A is diagonal, so is the corresponding 2-by-2 Cholesky factor L' such that $L'L'^T = A'$.

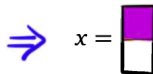
Let the vector  complete the 3-by-3 leading block $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ of A .


Then, the up-looking Cholesky algorithm requires that we do the sparse triangular solve of $L'x = b$.

The reach of node 1
in the graph
associated with L' is
just 1 itself.








- The Cholesky factor $L' := \begin{bmatrix} L' & 0 \\ x^T & w \end{bmatrix}$ of $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ has structure .

Up-looking Cholesky algorithm, cont'd₃



- Consider a matrix A with the following non-zero pattern:

- Let the vector $\begin{bmatrix} \text{blue} \\ \text{white} \\ \text{white} \\ \text{white} \end{bmatrix}$ complete the 4-by-4 leading block $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ of A .
 b

Then, the up-looking Cholesky algorithm requires that we do the sparse triangular solve of $L'x = b$.

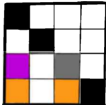
$$\begin{bmatrix} \text{black} & \text{white} & \text{white} \\ \text{white} & \text{black} & \text{white} \\ \text{purple} & \text{white} & \text{black} \end{bmatrix} \begin{bmatrix} \text{orange} \\ \text{orange} \\ \text{orange} \end{bmatrix} = \begin{bmatrix} \text{blue} \\ \text{white} \\ \text{white} \end{bmatrix}$$

$L' \quad x \quad b$


The reach of node 1
in the graph
associated with L' is 1
and 3.



$$\Rightarrow x = \begin{bmatrix} \text{orange} \\ \text{white} \\ \text{orange} \end{bmatrix}$$

- The Cholesky factor $L' := \begin{bmatrix} L' & 0 \\ x^T & w \end{bmatrix}$ of $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ has structure .


Up-looking Cholesky algorithm, cont'd₄





- Consider a matrix A with the following non-zero pattern:

- Let the vector b complete the decomposition $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ of A .


Then, the up-looking Cholesky algorithm requires that we do the sparse triangular solve of $L'x = b$.



The reach of nodes 1 and 2 in the graph associated with L' is 1, 2, 3, and 4.





- The Cholesky factor $L' := \begin{bmatrix} L' & 0 \\ x^T & w \end{bmatrix}$ of $\begin{bmatrix} A' & b \\ b^T & a \end{bmatrix}$ has structure





Up-looking Cholesky algorithm, cont'd₅

- ▶ The up-looking Cholesky algorithm yield a factor with the following non-zero pattern:



Note that the sparsity of L resembles that of A , with additional **fill-ins**:




- ▶ While the up-looking algorithm is better than performing a dense Cholesky factorization, it does **require many DFS** in graphs.
- ▶ We'll now try to do better than the up-looking algorithm.

Elimination tree


- ▶ The **graph associated with** the sparsity pattern of a **Cholesky factor** L has a special property which allows to **ignore many of its edges** and **retain the same reach**.
- ▶ Consider what happens when we **ignore all the non-zero entries of L below the first subdiagonal non-zero component**. E.g.,



Removing these entries results in a **sparsification of the associated graph**:



Elimination tree, cont'd



Two general properties are observed:


- 1 The **reach** of every vertex remains **unchanged** by sparsification.
- 2 **Every vertex** of the sparsified graph **has at most one edge** leading **out** of it.
I.e., **if** the graph is **connected**, then it is a **directed tree**.

Remarks:

- The sparsified graph is called an **elimination tree**.
- The elimination tree **may be disconnected**, in which case it is a forest, but even then, it will be called an elimination tree.
- The elimination tree is an important data structure that **can be used to simplify** all **reach calculations** in a sparse Cholesky factorization.

Non-zero pattern of L

- Say we aim to compute the i -th line of the Cholesky factor L of an SPD A .
- We are equipped with $L' := L[1:i-1, 1:i-1]$:



- We saw the non-zero entries of $L[i, 1:i-1]$ are the non-zero entries of the solution x of the above system with right-hand side $b := A[1:i-1, i]$.
- Remember from our sparse triangular solves, $x_j = \ell_{ij}$ is non-zero either if (a) $b_j = a_{ij} \neq 0$, or if (b) $\exists k < j$ so that both $\ell_{jk} \neq 0$ and $x_k = \ell_{ik} \neq 0$.

Non-zero pattern of L , cont'd

- Therefore, the pattern of non-zero values of L is characterized as follows:

Graph of (possible) non-zero entries of L

Let $j < i$, then ℓ_{ij} is non-zero if

- (a) $a_{ij} \neq 0$, or
- (b) there is a column index $k < j$ such that ℓ_{jk} and ℓ_{ik} are non-zero.

We denote by G_{ch} the graph with fewest edges that respect (a) and (b).

That is, G_{ch} is the minimal graph such that $a_{ij} \neq 0 \implies (j, i) \in G_{ch}$ and $(j, k), (i, k) \in G_{ch} \implies (j, i) \in G_{ch}$.

- The graph G_{ch} is a superset of the non-zero pattern of the Cholesky factor L of A .

It can be that $(j, i) \in G_{ch}$ but ℓ_{ij} numerically cancels out. However, if so, a tiny perturbation of A with fixed sparsity is enough to make $\ell_{ij} \neq 0$.

Therefore, the graph G_{ch} is best referred to as the graph of possible non-zero entries of the Cholesky factor L .

Definition of the elimination tree

► The elimination tree can be defined as follows:

Elimination tree

Let A be an SPD matrix, and G_{ch} be the graph representing the non-zero entries of the Cholesky factor L of A .


The elimination tree is obtained as follows.

For each node i in G_{ch} :

- Let V_i be the set of nodes j of G_{ch} for which there is an edge $(i, j) \in G_{ch}$, i.e., V_i is the set of out-neighbors of i . Let $p_i = \min V_i$ be the smallest-indexed node in V_i .
- Remove the edges (i, j) for all $j \in V_i \setminus \{p_i\}$ from G_{ch} , i.e., remove all the edges leaving i except for (i, p_i) .


The **elimination tree** is what's left of G_{ch} .

Example:



Properties of the elimination tree

- ▶ Since for each vertex i in G_{ch} , the elimination tree is formed by removing all but one out-neighbors, each vertex is left with at most one single out-neighbor, and the elimination tree is indeed a tree, or at least a forest.
- ▶ Consider the following example for a graph G_{ch} of non-zero entries of the Cholesky factor L :



As G_{ch} is, the reach of k is j, i .

If $k < j < i$, the elimination tree is formed by removing the edge (k, i) .

Then, the reach of k in the elimination tree is still j, i .

Theorem (Conservation of reach)

For a given graph G_{ch} of non-zero entries of a Cholesky factor L of A , for any $1 \leq i \leq n$, the reach of the corresponding elimination tree is the same as the reach of i in G_{ch} .

Computing the elimination tree from A

- ▶ Since the elimination tree has the same reach as G_{ch} , but is sparser than G_{ch} , it can be used to more efficiently identify the non-zero entries of the Cholesky factor.

For that, we need to figure out how to efficiently compute the elimination tree from the given sparsity pattern of A .


- ▶ The idea behind computing the elimination tree of A is to proceed one vertex at a time, maintaining a forest which contains all the vertices added so far. The elimination tree shall be obtained once all the vertices are added.
- ▶ Suppose we have a forest which has all the vertices $1, \dots, i-1$ at the correct place. To proceed with the i -th vertex, if $a_{ik} \neq 0$ for some $k < i$, then we'll want i to be in the reach of k . In order to avoid potential redundant edges, we should then connect i to whichever vertex j which is at the leaf of the tree containing k .

Computing the elimination tree from A , cont'd

► The pseudocode of the algorithm to build the elimination tree from the sparsity pattern of A is given by

1. Initialize a forest $\mathcal{F} = \emptyset$:
2. For $i = 1, \dots, n$:
3. Add vertex i to \mathcal{F}
4. For all $k < i$ such that $a_{ik} \neq 0$:
5. Find vertex j at the leaf of k 's tree
6. Add the edge (j, i) to \mathcal{F}

Taking the same sparse matrix A as earlier, the elimination tree is then built as follows:



We see that the same elimination tree is obtained as before.

Summary

To compute a sparse Cholesky factor L of a sparse matrix A , we

- 1 Build the elimination tree of A , at cost $\mathcal{O}(|A|)$, where $|A|$ is the number of non-zero entries in A .
- 2 Find the graph G_{ch} of possible non-zero entries of L using reaches of the elimination tree.
- 3 Perform the up-looking Cholesky factorization to build L .

Pseudocode of the up-looking Cholesky factorization to build row k of L :


1. $L[k, 1 : k] := A[k, 1 : k]$
2. For each $j < k$ such that $\ell_{kj} \neq 0$:
3. $\ell_{kj} \leftarrow \ell_{kj} / \ell_{jj}$
4. For each $i > j$ such that $\ell_{ij} \neq 0$:
5. $\ell_{ki} \leftarrow \ell_{ki} - \ell_{ij} / \ell_{kj}$

Nested dissection

Section 9.5 in Darve & Wotters (2021)

Reducing fill-ins in L

- While the row and column permutations of a matrix do not really impact the solution of a linear system (i.e., $P_r A P_c \cdot P_c^T x = P_r b$), they can have a significant impact on the sparsity pattern, i.e., the graph G_{ch} of the Cholesky factor:



Even though only one row and one column are permuted between A_1 and A_2 , the difference between the numbers of fill-ins in L_1 and L_2 is very significant.


- How should a matrix be ordered to reduce the number of fill-ins in the Cholesky factor L ?

One step of nested dissection

- ▶ **Nested dissection** is a strategy for ordering a matrix A in a way that **closely minimizes** the number of **fill-ins** in L .
- ▶ **Nested dissection** is a **recursive method based on graph partitioning**.
- ▶ Consider the symmetric matrix A with an associated graph G .

Let the vertices of G be decomposed in the disjoint union of V_1, V_2 and S , so that there are no edges between vertices of V_1 and V_2 .

If G_1 and G_2 are the induced graph on V_1 and V_2 , respectively, then we have




where S is referred to as a **separator**.

One step of nested dissection, cont'd₁

- ▶ A **node separator set** S partitions the graph G of A into three disjoint sets of vertices V_1, V_2 and S such that none of the nodes of V_1 are connected to any of the nodes of V_2 , and vice-versa.
- ▶ The removal of S from the graph G leads to two subgraphs G_1 and G_2 , disconnected from each other.
- ▶ Consider what happens when we order the vertices as


(vertices of G_1 , vertices of G_2 , S)

we obtain the matrix



which is structurally close to the A_1 matrix with small number of fill-ins.

One step of nested dissection, cont'd₂



As a result of the block diagonal structure due to G_1 and G_2 , the Cholesky factor L of the reordered matrix will preserve a block diagonal structure. If the blue blocks are dense, the sparsity of L is exactly given by that of A . In general, each block of L will be sparse.

From here, we will proceed as follows:

- ▶ **Entries in S .** For these entries, we give up and accept whatever fill-ins happen. Thus we want S to be as small as possible.
- ▶ **Entries in G_1 and G_2 .** For these entries, we will recurse on the blocks $G_1 \times G_1$ and $G_2 \times G_2$, i.e., find small separators S_1 and S_2 for G_1 and G_2 , respectively, and so on.

Nested dissection

- ▶ The basic idea of **nested dissection** is to **recursively** apply the procedure we just described, and yield a **nested dissection ordering** of the graph nodes.


Pseudocode for the nested dissection algorithm

```
function nested_dissection(G::Graph)
    # returns a nested dissection ordering of the nodes of G
    if size(G) < threshold
        # This is the base case of the recursion
        return nodes(G) # nodes of graph G
    end
    [G1, G2, S] = find_separator_set(G)
    # G1 and G2 are two disconnected subgraphs
    # S is the node separator set
    P1 = nested_dissection(G1)
    P2 = nested_dissection(G2)
    return [P1; P2; S]
end
```

- ▶ The description of `find_separator_set` is beyond the scope of this class. A good separator set has as few nodes as possible, and it decomposes the graph in roughly equally sized subgraphs. Finding a good separator set is actually a NP-hard problem.

Nested dissection, cont'd

- In the matrix, the recursive process of nested dissection looks like this:



As we can see, if good separators are chosen, the "down-and-right-arrow" patterns shows up at all scales, and we can guarantee that more and more entries of L will be zero.

Homework problems

Homework problem

Turn in **your own** solution to the following problem:

Pb. 16 Find the non-zero pattern of the Cholesky factor L for the following matrix:

$$A = \begin{bmatrix} a_{11} & a_{12} & 0 & a_{14} \\ a_{12} & a_{22} & 0 & a_{24} \\ 0 & 0 & a_{33} & 0 \\ a_{14} & a_{24} & 0 & a_{44} \end{bmatrix}.$$

Show your work using the up-looking Cholesky factorization algorithm.